



Running Dataproc jobs

Data Engineering on Google Cloud Platform



©Google Inc. or its affiliates. All rights reserved. Do not distribute.
May only be taught by Google Cloud Platform Authorized Trainers.

Notes:

25 slides + 1 lab: 1 hour

Cloud Dataproc provides compelling reasons to run open-source tools on GCP

- Stateless clusters in <90 seconds *MODULE 1*
- Supports Hadoop, Spark, Pig, Hive, etc.
- High-level APIs for job submission
- Connectors to Bigtable, BigQuery, Cloud Storage



Notes:

We have already looked at #1.

Let's look at #2 and #3 here. Starting with #2.

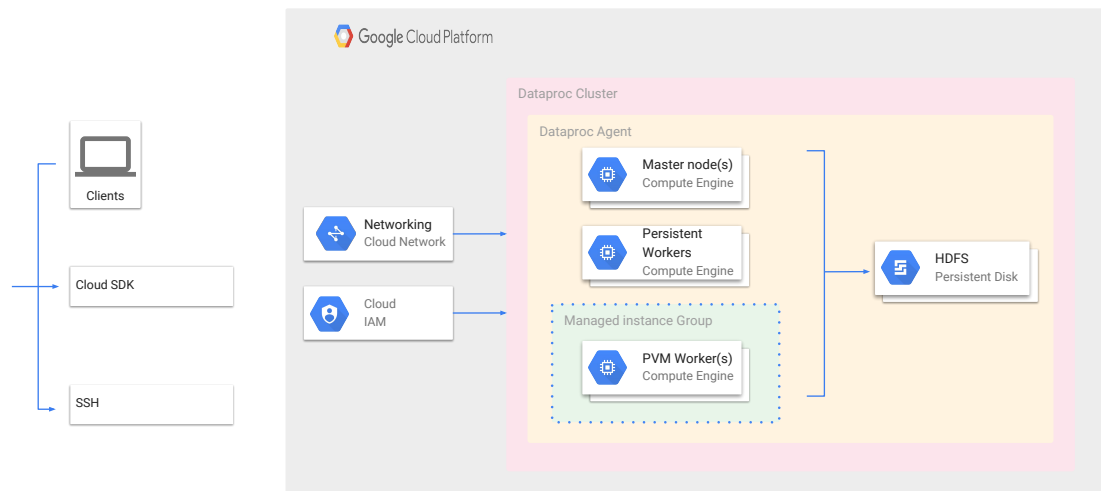
Agenda

Running jobs + Lab

Separation of storage and compute

Submitting jobs + Lab

Can SSH to cluster and run Pig/Spark



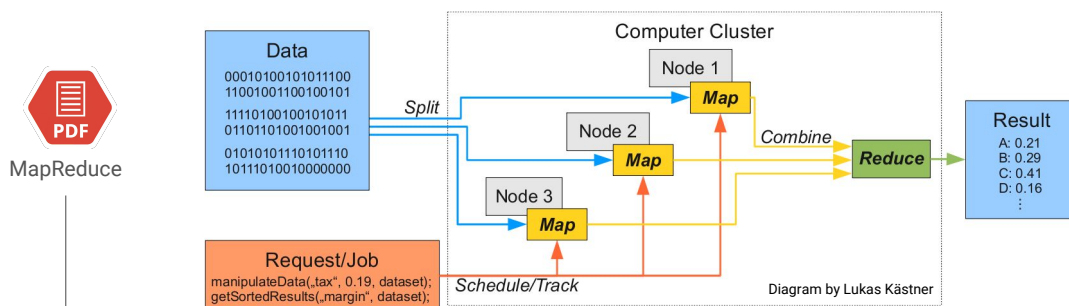
Lab - Leveraging Unstructured Data : Part 2

- SSH into the cluster to run Pig and Spark jobs interactively
- Work with HDFS

Agenda

Separation of storage and compute

MapReduce approach splits Big Data so that each compute node processes data local to it

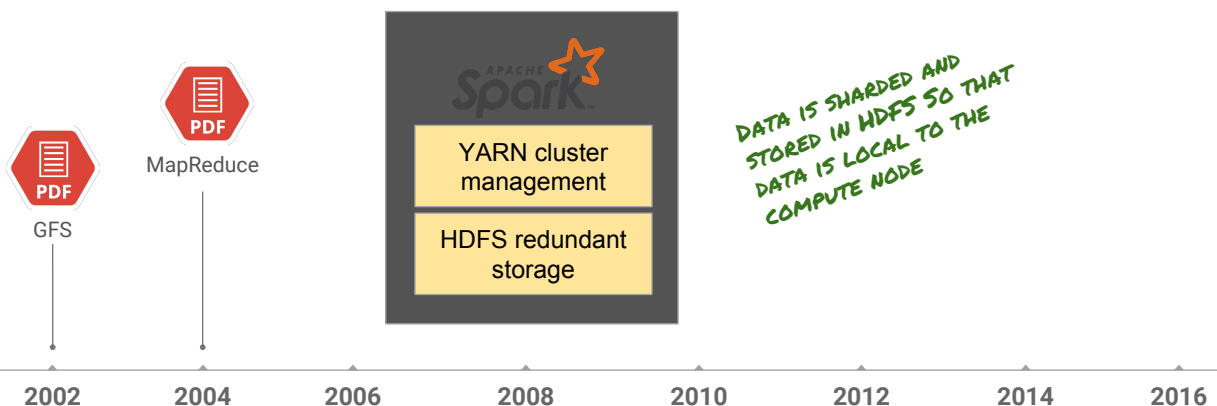


Notes:

This slide was also in Chapter 1, so just mention that they've already seen this.

Diagram source: <https://www.flickr.com/photos/lkaestner/4861146813>
cc-by-sa Lukas Kastner

To get data local to the machine, you pre-shard the data onto Hadoop Distributed File System



Notes:

HDFS is based on the 2002 paper from Google on Google File System.

Compute and Storage are closely tied in traditional MapReduce architecture

Scenario	What needs to happen?
Compute node needs to be replaced	?
Append new year of data	?
?	?
?	?

Notes:

Ask the class what problems this leads to. Have them think about a scenario.

Example scenario: You split your data into 10 nodes. One node goes down. You bring in a new replacement. What has to happen? At least some part of the data needs to be copied onto the new nodes before jobs have to be partitioned.

Example scenario: The data changes (maybe you need to change formats or append a new year of data). What needs to happen?

Need to provision increased resources quickly



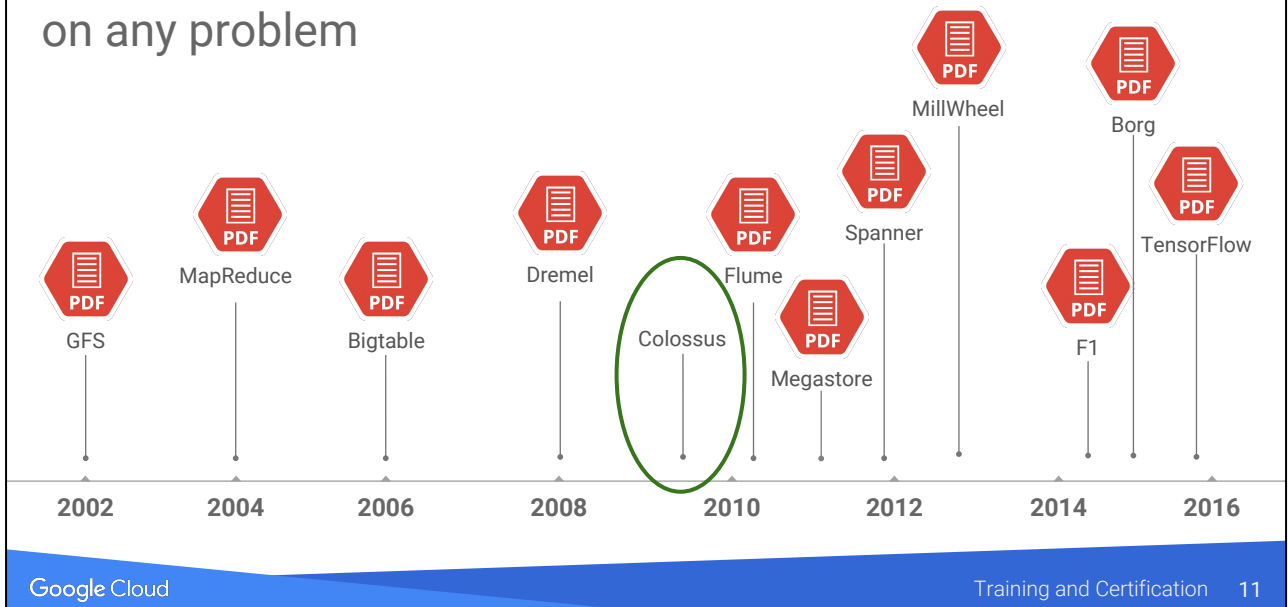
Notes:

Beyond the obvious question of deciding how much to provision, the larger issue is that of the risk/cost of experimentation.

If they knew for sure what they have to do, how to do it, and what resources they need for it, enterprises would be able to provision large amounts of resources. But the reality is that you have to experiment. It's not about how much you can invest, it's how quickly you can iterate.

Image source: I (vbp@) took it myself.

Bring the power of the datacenter to bear on any problem

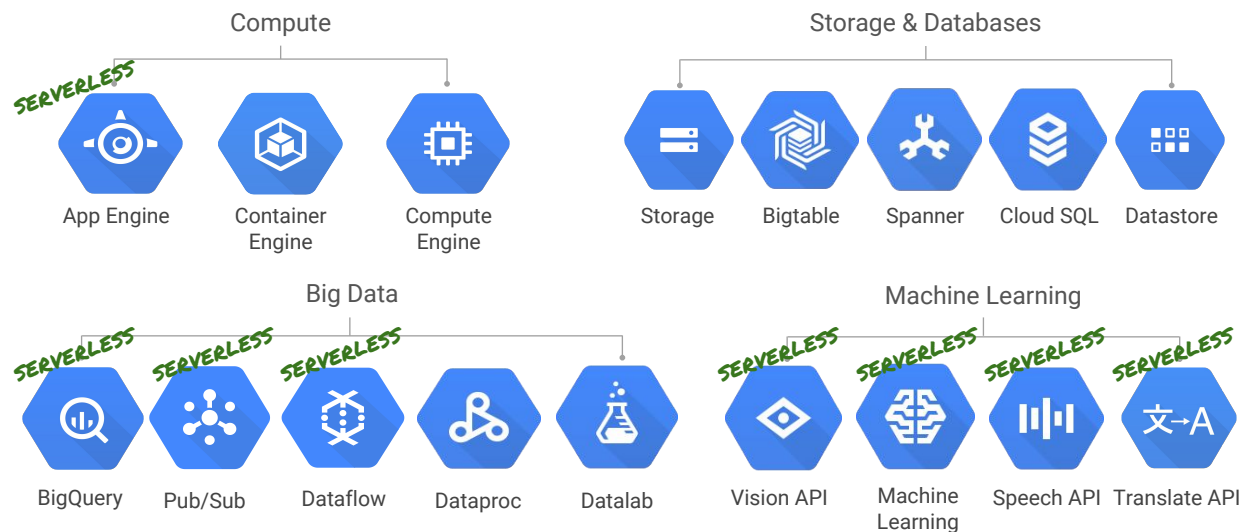


Notes:

Colossus, the replacement for GFS, is the key innovation and led to a bunch of serverless offerings. It's in our datacenter and enables you to not have to shard data. Instead, you have a global filesystem that offers petabit/second bisection bandwidth. Yes, Colossus has not been published ... public information on it is scarce and consists of an unofficial copy of a slide deck by Andrew Fikes:

<https://www.systutorials.com/3306/storage-architecture-and-challenges/>

GCP gives you access to that power



Notes:

The Big Data and ML offerings are serverless (except for Dataproc & Datalab because they are based on OSS -- Hadoop ecosystem and Jupyter respectively). The APIs are of course serverless although you tend not to think of them as serverless offerings.

GCP gives you serverless platform for all stages of the analytics data lifecycle

Ingest



Pub/Sub

Processing



Dataflow

Analysis



BigQuery

*NO NEED TO GUESS
CAPACITY OR WORRY
ABOUT IDLE RESOURCES
NOTHING TO MAINTAIN*

Notes:

In particular, the three Big Data serverless offerings are ...

Serverless data processing is about speed, low cost, and freedom

Speed to insights

Focus on insights

Not administration

Low cost

Practically infinite scale, exactly when you need it

Pay only for what you use

Freedom to experiment


Experiment, fail quickly, and iterate

Successful experiments are ready to go live right away

Notes:

Summary slide for this section: Not just about low-cost, but also about speed and freedom.

Separation of Storage and Compute is what enables Serverless to work




Notes:

Separation of storage and compute is what enables “serverless to work” -- Dataflow can read use any of these as source/sink. This sort of direct read is efficient because of very high sustained read speed from Cloud Storage -- any two computers in data center are connected by very fast network.

Keep as much data as you want, economically.

Share data in place, no more FTP and copying.

BigQuery also separates compute and storage




Notes:

Access any storage system from any processing tool.

Compare and contrast bq's data separation w/ typical DB storage mgmt system.

BigQuery is "just" a query engine. It can query csv files on cloud storage also, for example.

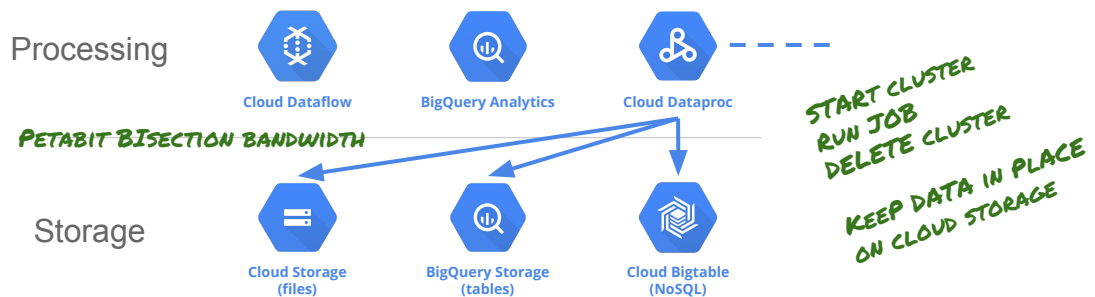
Can I run Spark and still get separation of storage and compute?



Notes:

But what if I want to run Spark programs? How can I get separation of compute & storage. Spark runs on Hadoop ... and Hadoop is a cluster-aware piece of software ... we need to take our data and split it into pieces and store them on cluster so that data is local to compute ... but then we are limited by the number of processing nodes or the number of storage nodes. These are not independent

Cloud Dataproc provides the ability for Spark programs to also separate compute & storage



Notes:

Just change all your input urls from `hdfs://` to `gs://` ...


The reason you can do this is the speed of the inter-networking

Cloud Dataproc is Spark/Hadoop "the Cloud way"

Deploy cluster in ~90 seconds

Pay by the minute

Sustained reads from Cloud Storage are even faster than before ...



Notes:

The impact of the PB/s networking and software improvements is that sustained reads are very fast. You can keep things in GCS.

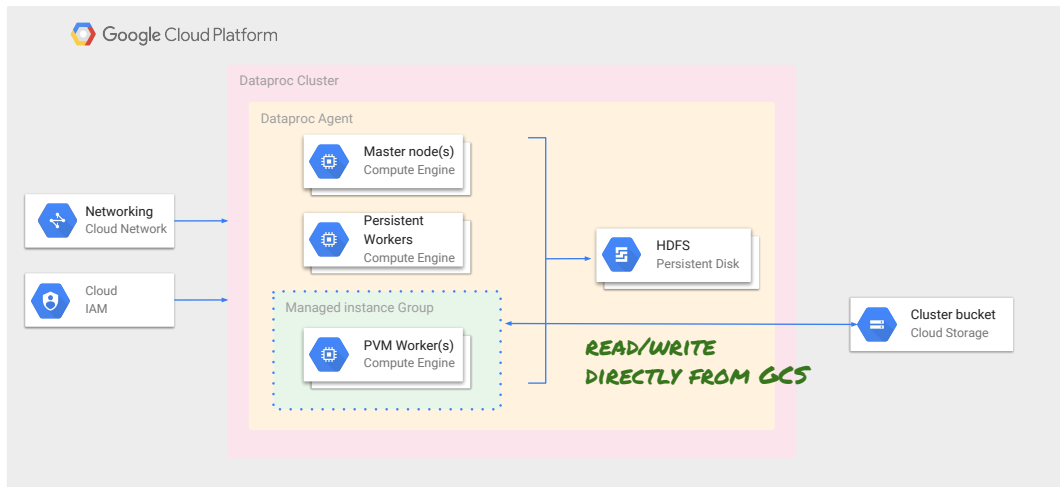
Old: 2016

New: 2017

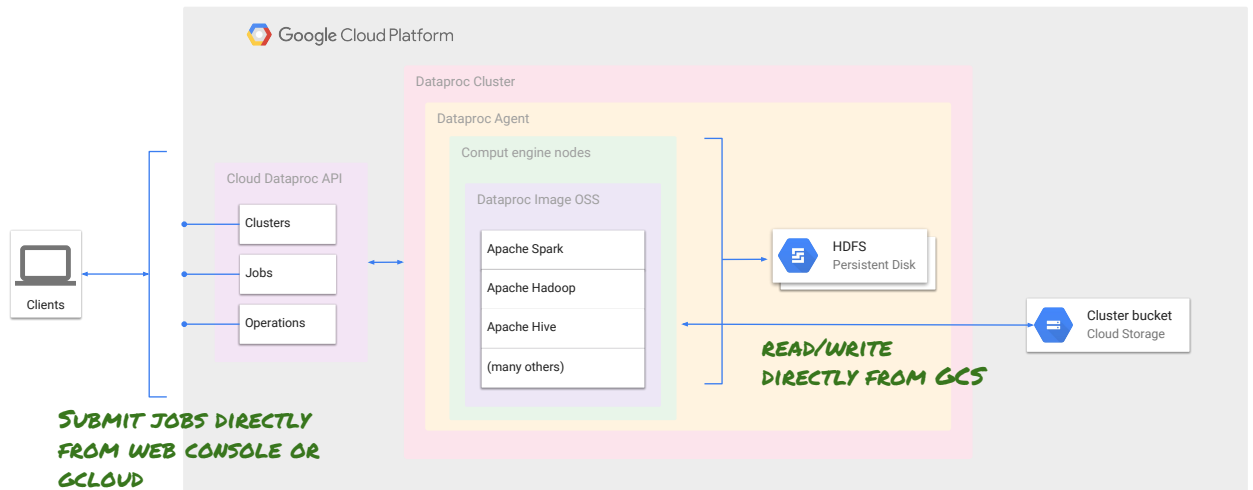
Agenda

Submitting jobs + Lab

Cloud Dataproc hardware architecture



Cloud Dataproc software architecture



Notes:

Can directly read/write to GCS from Spark, Pig, etc.
Submit jobs

Lift and shift work to Cloud Dataproc

1

Copy data to GCS

Copy your data to Google Cloud Storage (GCS) by installing the connector or by copying manually

2

Update file prefix

Update the file location prefix in your scripts from `hdfs://` to `gs://` to access your data in GCS

3

Use Cloud Dataproc

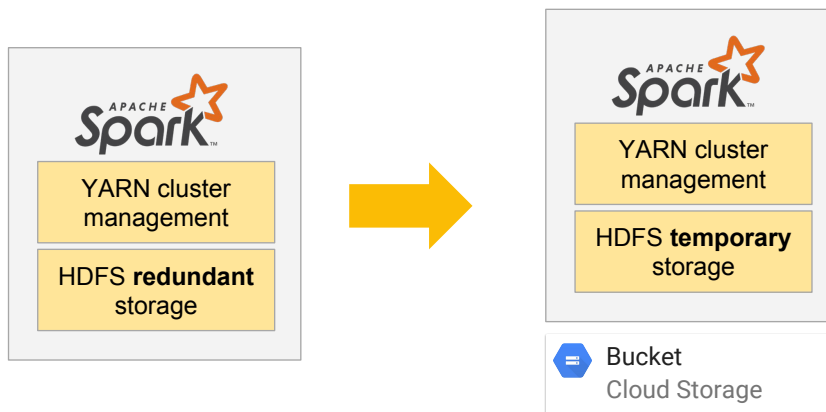
Create a Cloud Dataproc cluster and run your job on the cluster against the data you copied to GCS. Done

Migrating code

- In most cases, you only need to update jobs so they read from Google Cloud Storage (gs://) instead of HDFS

```
textFile = sc.textFile("hdfs://...") # Read data  
  
# Creates a DataFrame having a single column  
df = textFile.map(lambda r: Row(r)).toDF(["line"])  
errors = df.filter(col("line").like("%ERROR%"))  
# Counts all the errors  
errors.count()  
# Counts errors mentioning MySQL  
errors.filter(col("line").like("%MySQL%")).count()  
# Fetches the MySQL errors as an array of strings  
errors.filter(col("line").like("%MySQL%")).collect()
```


Why change hdfs:// to gs://?



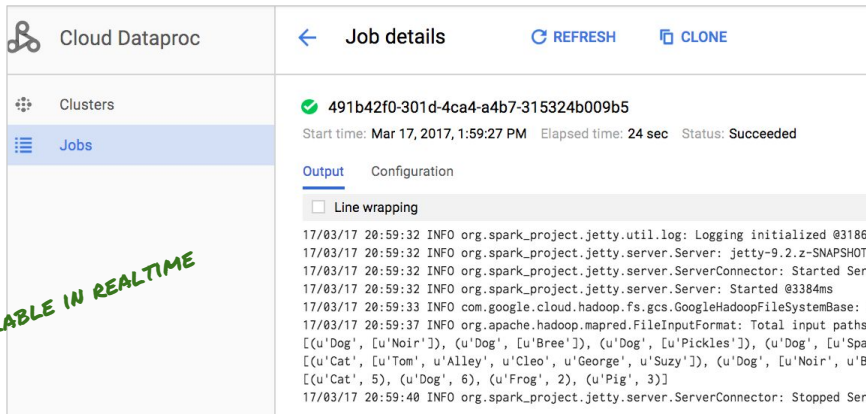
Notes:

Because the cluster is temporary We want to be able to delete the cluster when we are done.

If the first case, HDFS is the durable storage for data. We can't delete the cluster.

In the second case, HDFS is only temporary. We can delete the cluster.

Monitor logs of submitted jobs from web console



Cloud Dataproc

← **Job details** [REFRESH](#) [CLONE](#)

✓ **491b42f0-301d-4ca4-a4b7-315324b009b5**

Start time: Mar 17, 2017, 1:59:27 PM Elapsed time: 24 sec Status: Succeeded

Output Configuration

☐ Line wrapping

```
17/03/17 20:59:32 INFO org.spark_project.jetty.util.log: Logging initialized @3186
17/03/17 20:59:32 INFO org.spark_project.jetty.server.Server: jetty-9.2.z-SNAPSHOT
17/03/17 20:59:32 INFO org.spark_project.jetty.server.ServerConnector: Started Ser
17/03/17 20:59:32 INFO org.spark_project.jetty.server.Server: Started @3384ms
17/03/17 20:59:33 INFO com.google.cloud.hadoop.fs.gcs.GoogleHadoopFileSystemBase: (
17/03/17 20:59:37 INFO org.apache.hadoop.mapred.FileInputFormat: Total input paths
[(u'Dog', [u'Noir']), (u'Dog', [u'Bree']), (u'Dog', [u'Pickles']), (u'Dog', [u'Spa
[(u'Cat', [u'Tom', u'Alley', u'Cleo', u'George', u'Suzy']), (u'Dog', [u'Noir', u'Bi
[(u'Cat', 5), (u'Dog', 6), (u'Frog', 2), (u'Pig', 3)]
17/03/17 20:59:40 INFO org.spark_project.jetty.server.ServerConnector: Stopped Ser
```

Notes:

These logs are available in real-time.

Monitor cluster usage graphs on Web UI, Stackdriver



Notes:

Overall cluster usage from Dataproc page.

Individual VMs from Compute Engine VM.

Lab - Leveraging Unstructured Data : Part 3

- Create a Cloud Storage bucket to store job input, output, and application files
- Submit jobs using the Web Console
- Submit jobs using the CLI
- Monitor job progress and view results



cloud.google.com